

Persuasion in Scene: A Multi-Agent Typographic Jailbreak Crew against Large Vision-Language Models

Anonymous CVPR submission

Paper ID 21049

Abstract

Large Vision-Language Models (LVLMs) have achieved remarkable performance in various domains. However, their integration of new visual modality into base Large Language Models (LLMs) may introduce additional security vulnerabilities. To explore the safety boundaries of LVLMs which contribute to building more trustworthy and reliable models, jailbreaking LVLMs has been a preliminary and essential research direction. Existing jailbreaking methods are facing increasing challenges with the development of visual security detection mechanism. Moreover, these approaches overemphasize images design while neglecting the role of textual prompts, resulting in weak cross-modal synergy. To address both limitations, we propose a novel framework named *Persuasion in Scene (PIS)*: a multi-agent typographic jailbreak crew against LVLMs. The crew comprise three specialized agent teams: *PROMPTER*, *PAINTER*, and *GUIDER*. Each team is equipped with a supervisor to ensure its work efficiency and quality. Extensive experiments demonstrate the effectiveness of our proposed PIS. PIS achieves an average attack success rate (ASR) of over 60% against advanced LVLMs including GPT-4o, Gemini 2.5 Flash, Qwen3-VL-Plus, and GLM-4.5V, significantly outperforming other challenging methods.

1. Introduction

Large Vision-Language Models (LVLMs) have recently demonstrated remarkable capabilities in understanding multimodal information, enabling various applications such as image captioning and visual question answering. Due to outstanding performance in visual-textual content alignment and processing of multimodal models such as GPT-4o [12], Gemini [25], Qwen-VL [1], and LLaVA [15], they are widely used in healthcare diagnostics [10, 30] and autonomous driving [6, 27, 34]. However, this rapid proliferation has simultaneously raised public concerns about safety. Although LVLMs inherit safety alignments from their foun-




Figure 1. A comparison of three scenarios reveals differential GPT-4o results: Case 1 (fixed prompt + typographic image), Case 2 (simple prompt + concatenated image), and Ours (structured persuasive prompt + scenario typographic image).

dational Large Language Models (LLMs), recent studies [8, 13, 21, 23] have revealed that their increased vulnerability to adversarial attacks, largely attributable to the introduction of image modality. Consequently, investigating jailbreak in LVLMs is imperative research for securing multimodal systems.

To exploit the vulnerability of visual safety mechanisms in LVLMs, existing studies [9, 21, 31] propose white-box attack methods using model gradient information to generate adversarial images for jailbreaking. In contrast, black-box attacks for closed-source models assume attackers can only modify inputs for harmful content generation. A key paradigm in black-box attacks is the use of “typography” technology. For example, FigStep [8] typesets text instructions onto images indirectly to achieve the transfer of harmful information, while other methods [13, 16, 17] “concatenate” typographic images with auxiliary visuals such as scene-related images to steer the target model toward generating harmful responses. Although these methods have

proven effective for early LVLMs, their limitations have gradually emerged due to the improved visual detection tools for harmful information in these attack images and advances in cross-modal security mechanisms. Through systematic analysis, we empirically demonstrate two key limitations in current black-box attack paradigms using “typography and concatenation” techniques. First, the attack images constructed by these methods lack continuity and naturalness, making the harmful information so distinct that advanced LVLMs can detect it easily and trigger the security mechanism, leading to attack failure (Fig. 1). Second, existing methods overemphasize image design while neglecting textual prompts design, causing an imbalance in cross-modal attacks. However, recent studies [17, 24] have shown that cross-modal consistency between textual prompts and images is crucial for successful LVLm jailbreaking. To address both limitations, we hypothesize and experimentally verify that elementizing harmful information and integrating it naturally into scenario-relevant images can bypass visual security detection. Moreover, we use logical rewriting and data citation strategies to design “structured persuasive prompts”, which enable the text modality to participate in jailbreaking actively.

Based on the above analysis, we propose a black-box framework **Persuasion in Scene (PIS)**, which is a novel multi-agent crew for jailbreaking against LVLMs. PIS automatically generates a structured persuasive prompt and a scenario typographic image from an original harmful instruction, with the entire process driven by agent teams without human intervention. Specifically, the crew operates in three collaborative team: the **PROMPTER** team rewrites the harmful instruction using logical rewriting and data citation strategies; The **PAINTER** team embeds extracted harmful keywords into a scene image as reasonable clue, thereby ensuring the image is natural and continuous; and the **GUIDER** team sanitizes the text and enhances cross-modal synergy between the generated text and image.

In summary, our key contributions are as follows:

- We systematically analyze the limitations of existing black-box attacks based on “typography and concatenation” techniques, identifying two issues: lack of image naturalness and imbalance in cross-modal coordination.
- We introduce PIS, a novel jailbreak framework leveraging multi-agent crew collaboration. The framework comprise three specialized agent teams (PROMPTER, PAINTER, GUIDER) overseen by a supervisory mechanism. This enables PIS to achieve end-to-end, automated, and training-free attack samples generation while ensuring high cross-modal consistency between generated “structured persuasive prompts” and “scenario typographic images”.
- Extensive experiments demonstrate that PIS achieves an average Attack Success Rate (ASR) exceeding 60% across various mainstream closed-source and open-source

LVLMs, including GPT-4o, Gemini 2.5 Flash, Qwen3-VL-Plus, and GLM-4.5V. This performance significantly surpasses other challenging methods, validating the effectiveness and robustness of our framework.

2. Related Work

2.1. Jailbreak Attacks against LLMs

In recent years, Large Language Models (LLMs) have been widely adopted in practical applications. Meanwhile, research directions focused on jailbreaking them have grown more diverse. These jailbreaking methods can be roughly categorized into two types: gradient optimization-based approaches represented by GCG [36] and gradient-free prompt engineering methods represented by “Do-Anything-Now (DAN)” series [28]. Notably, among gradient-free methods, a paradigm focused on generating persuasive prompts through strategies such as logical appeal, role-playing, and scenario construction has recently emerged as a prominent trend. For instance, PAP [32] proposes a taxonomy of persuasion strategies for automatic prompt generation; PAIR [2] iteratively refines prompts through “attacker-target model” interaction; Shah et al. [22] leverage persona design to elicit harmful outputs from target model. Despite significant progress in LLMs jailbreaking, this persuasive logic paradigm is confined to the textual modality due to LLMs’ unimodal nature. We not only apply such logic to textual prompts rewriting but also extend it to the visual modality such as generating immersive inducing images through role-playing and scenario-based design.

2.2. Jailbreak Attacks against LVLMs

Based on different access rights to target models, jailbreak attacks on LVLMs are categorized into white-box and black-box attacks. White-box methods [9, 21, 31] rely on model parameters and gradient information. However, black-box attacks do not require internal model information, instead achieving jailbreak by transferring harmful textual information to the visual modality. For instance, FigStep [8] embeds harmful text onto blank images using typographic technology; Some studies [13, 16] find that the relevance between input images and harmful information is positively correlated with attack success rate. Therefore, they generate query-relevant images and concatenating them with typographic images containing harmful information to make an attack image; Moreover, SI-Attack [33] enhances attack effectiveness by shuffling attack images and textual prompts. However, with the improvement of LVLMs’ security mechanisms, discontinuous images generated by typography and concatenation techniques are prone to visual security detection of hidden harmful information, leading to model refusal [18]. Moreover, recent studies [17, 24] indicate that text-image cross-modal consis-

tency is the key to successful jailbreaking, but most existing methods overemphasize image design while underemphasizing the design of text, resulting in imbalanced attacks.

2.3. Evolution of Automation and Agent Frameworks for Jailbreak Attacks

Early jailbreak attacks relied heavily on manual operations, suffering from low efficiency, poor scalability, and difficulty covering complex scenarios. With technological evolution, automated jailbreak sample generation has become a trend. For LLMs jailbreaking, PAIR [2] and PAP [32] automatically generate jailbreak prompts for target black-box LLMs by constructing an attacker LLM; For LVLMs jailbreaking, studies like FigStep [8], HADES [13], HIMRD [17], have also proposed methods to automatically generate jailbreak prompts and images. Meanwhile, multi-agent systems have emerged as a key direction for solving complex tasks due to their strong collaboration. AgentVerse [4] simulates human group problem-solving processes, demonstrating excellent performance in some tasks; The AutoAgents [3] framework adaptively generates and coordinates multi-agents based on task requirements, with an additional observer role to optimize architectural division of labor. These studies provide an architectural reference for designing framework which can automatically generate jailbreak samples in this paper.

3. Methodology

3.1. Preliminaries

Jailbreak Attack Against LVLMs. We define a LVLM as $\mathcal{M}_\theta$, with the text and vision domains denoted as $\mathbb{T}$ and $\mathbb{V}$, respectively. The model receives a text input $x_t \in \mathbb{T}$ and a visual input $x_v \in \mathbb{V}$ to generate a response $y \in \mathbb{T}$. A jailbreak attack on the target model $\mathcal{M}_\theta$ aims to elicit a desired harmful output $y_t \in \mathbb{T}$. This requires the attacker to design a jailbreak strategy. Common strategies include exploiting the weaker security mechanisms of the visual modality by transfer the harmful information $t \in \mathbb{T}$ into the visual input x_v via a strategy ϕ_v . Alternatively, the harmful information is distributed across both modalities via strategies ϕ_t and ϕ_v to generate x_t and x_v , thereby circumventing the LVLM’s multimodal security defenses. The objective of a jailbreak attack on an LVLM can be expressed as:

$$\underset{\phi_t, \phi_v}{\text{Minimize}} \quad \mathcal{L}(y_t, \mathcal{M}_\theta(x_t, x_v)), \quad (1)$$

where $(x_t, x_v) = (\phi_t(t), \phi_v(t))$.

3.2. Motivation

As shown in Eq. (1), successful jailbreaking depends critically on designing an effective jailbreak strategy $\phi(\cdot)$. Existing black-box methods (e.g., typography and concatenation) sanitize text prompts, but harmful information in their

blunt visual presentation remains easily detectable by the visual security mechanisms in LVLMs, often causing attack failure [18]. To address this, **our first motivation is to evade advanced visual security detection.** We embed harmful content into the scene image instead of displaying it explicitly, creating a “scenario typographic image”. This transforms the harmful information into a contextual image clue, making the visual safety mechanism perceive it as a benign part of the image. Moreover, most existing methods prioritize the image modality and lack an effective text-modality strategy ϕ_t that can synergize with the image modality, limiting cross-modal collaboration. **Our second motivation is to enhance this guidance capability from the text modality.** We design a “structured persuasive prompt” to steer the target model toward interpreting the embedded clue within the image. To implement these strategies, we design a multi-agent collaborative framework that operates in an automated and train-free manner.

3.3. Persuasion in Scene: A Multi-Agent Typographic Jailbreak Crew

In this section, we introduce PIS’s three core agent teams: PROMPTER, PAINTER, GUIDER. Specifically, the PROMPTER and GUIDER teams collaborate to generate “structured persuasive prompts” that are benign and persuasive, and the PAINTER team constructs “scenario typographic images” that plausibly embed the harmful information within an image context. PIS achieves a highly effective, synergistic attack by transferring the malicious payload to the image and rendering the text prompt harmless.

PROMPTER: Crafting Structured Persuasive Prompts.

Existing jailbreaking methods against LVLMs often overemphasize the image modality, failing to exploit the full jailbreaking potential of text.

Persuasive Logic3: [Goal-Oriented]

Definition: Frames the instruction as the necessary initial step required to initiate a larger, ultimately justicial ultimate goal.

Figure 2. **An illustration of the Goal-Oriented logic from corpus M_l .** The other four logics include Propositional, Justified Negation, Priming, and Time Pressure.

To address this limitation, we design the PROMPTER agent team whose objective is to refine the original harmful instruction t into a “structured persuasive prompt”. To achieve this, we first manually curate a corpus M_l containing five distinct persuasive logics. A persuasive logic example is illustrated in Fig. 2. The *Logic Decider* $\mathcal{D}$ and *Logic Rewriter* $\mathcal{E}$ select the most suitable logic from M_l to rewrite t , denoted as $T_l = \mathcal{E}(\mathcal{D}(M_l), t)$. To bolster persuasiveness, the *Data Citer and Structurer* $\mathcal{F}$ enriches T_l by incorporating relevant data, yielding $T_f = \mathcal{F}(T_l)$. We




Figure 3. Overview of the Persuasion in Scene (PIS) framework, consisting of three specialized agent teams: PROMPTER generates structured persuasive prompts using logic rewriting and data citation; PAINTER embeds harmful keywords into scene images to create natural scenario typographic images; and GUIDER team sanitizes the text prompts and enhances cross-modal synergy.

observe that simple fusion can cause semantic drift. To mitigate this, $\mathcal{F}$ also receives the original instruction t , tasking it not only with data citation but also with structuring the final prompt. This structure highlights a instruction T' consistent with the original intent and appends an encourage suffix E to elicit a response. The prompt thus implicitly combines: *Data Citation + Logic Reformulation + Highlighted Instruction + Encourage Suffix*. This process is expressed as $T_p = \mathcal{F}(\mathcal{E}(\mathcal{D}(M_l), t))$, which simplifies to $T_p = \mathcal{F}(T_f) + T' + E$. A *Supervisor* agent $\mathcal{J}_t$ facilitates this process through an iterative loop, providing scores S_t and feedback F_t to $\mathcal{D}$, $\mathcal{E}$, and $\mathcal{F}$. The team refines the prompt until $\mathcal{J}_t$ is satisfied or a maximum iteration count is reached, outputting the optimized prompt $P_{t,best}$. The entire PROMPTER team operation is thus:

$$T_p = P_{t,best} = \mathcal{J}_t(\mathcal{D}, \mathcal{E}, \mathcal{F} \mid t, M_l). \quad (2)$$

PAINTER: Making Scenario Typographic Images.

The core task of this team is to embed harmful information into a contextually relevant image, making it a plausible clue for the model's interpretation. First, the *Keyword Extractor* $\mathcal{A}$ identifies a set of keywords $k = \mathcal{A}(t)$ from

instruction t . The *Scene Scripter* $\mathcal{B}$ then generates a script based on t , defining three elements: role (r), scene (s), and action (a), denoted as $(r, s, a) = \mathcal{B}(t)$. A critical challenge is to plausibly integrate the typographic information. We define a requirement $U(k)$: **the image must feature a natural information carrier (e.g., a whiteboard, screen) upon which the keywords (k) are artfully and naturally distributed**. The *Art Director* $\mathcal{C}$ fuses the script elements (r, s, a) with this requirement $U(k)$ to formulate a detailed, natural prompt for the image generation model $\mathcal{Q}$. This is described as $I = \mathcal{Q}(\mathcal{C}((r, s, a) \oplus U(k)))$. Similar to the PROMPTER team, a *Supervisor* ($\mathcal{J}_i$) oversees an iterative optimization process to produce the final image generation prompt, $P_{i,best}$. The PAINTER team's workflow is:

$$I = \mathcal{Q}(P_{i,best}) = \mathcal{Q}(\mathcal{J}_i)(\mathcal{C}((r, s, a) \oplus U(k))). \quad (3)$$

GUIDER: Text Sanitization and Alignment.

Although the persuasive prompt T_p from the PROMPT team is persuasive, it still contains explicit harmful content detectable by safety mechanisms. Therefore, we designed the GUIDER team to sanitize this text and to guide the target model to use the image to reconstruct the full instruction.

The *Text Expert* $\mathcal{G}$ receives T_p and the image I . Referencing the keywords k from PAINTER and a predefined benignification method library M_h , which includes techniques such as Variable Substitution and Keyword Scrambling (e.g., obfuscating a term as “10T-N3twork” as shown in Fig. 3), $\mathcal{G}$ neutralizes T_p by applying methods from M_h to semantic elements related to k , resulting in T_h . After this harmful-to-benign transfer, $\mathcal{G}$ enhances synergy by appending a *Note* suffix (as shown in Fig. 3) to T_h . This suffix, while maintaining the prompt’s structure, subtly directs the model to interpret the image, retrieve k , reconstruct the original intent, and respond via itemized points. A *Supervisor* $\mathcal{J}_h$ assists $\mathcal{G}$ in refining this final text, T_f . The operation of the GUIDER team is formalized as:

$$T_f = \mathcal{J}_h(\mathcal{G}(T_p, I | k, M_h)). \quad (4)$$

Our three-agent framework characterized by a clear division of labor, agent autonomy, and collaboration yields the final input $(x_t, x_v) = (T_f, I)$. The overall process generating an attack example is shown in Fig. 3. The judging criteria for the three *Supervisors* are detailed in Sec. 4.4. They return a binary judgment (e.g., satisfactory/unsatisfactory); if unsatisfactory, they also provide revision suggestions to specified agent roles. We provide a detailed walkthrough of an end-to-end example in the Supplementary Material, illustrating each stage of our framework and its associated cost. The material also contains the complete algorithm with iterative optimization and all system prompts used.

4. Experiment

4.1. Experimental Setup

Datasets. To validate our proposed PIS framework, we conduct experiments on two benchmark datasets: SafeBench from FigStep [8] and the dataset from HADES [13]. For SafeBench, we select its top 7 high-harm categories: Illegal Activities, Hate Speech, Malware Generation, Physical Harm, Fraud, Pornographic Content, and Privacy Invasion. Each category contains 50 harmful instructions, resulting in 350 samples. The HADES dataset covers 5 scenarios (Violence, Finance, Privacy, Self-Harm, and Animals), comprising 750 harmful instructions in total.

Target Models. We evaluate approaches against 7 representative LVLMS. These include four commercial closed-source models: GPT-4V [18], GPT-4o [12], the recently released Gemini 2.5 Flash [5] and Qwen3-VL-Plus [26]. We also conduct jailbreak experiments on three other popular open-source models: Qwen3-VL-8B [26], InternVL3-9B [35] and GLM-4.5V [11]. It should be noted that the scale of the GLM-4.5V we used is 106B, and due to its large scale which complicates local deployment, we utilize its API.

Comparison Methods. To ensure a fair comparison, all attack methods are evaluated under a strict black-box, single-turn interaction setting, using their recommended configurations. We compare PIS with four state-of-the-art (SOTA) jailbreak attacks: FigStep [8], MM-SafetyBench [16], HADES [13], and SI-Attack [33], as well as a Vanilla baseline. Additionally, FigStep_{pro} [8] and HIMRD [17], which employ a strategy of dispersing harmful information, are also added as supplementary baselines. We observed that these latter two methods exhibit notable performance differences between open-source and closed-source models. Therefore, we conduct an analysis to highlight the superior universality of PIS.

Evaluation Metrics. We adopt Attack Success Rate (ASR) and Toxic Score as harm evaluation metrics. For evaluation, we utilize Deepseek-V3 [7] as the judging model. This model was selected based on experimental comparisons and manual analysis, which demonstrated that its scoring provides high discriminative power while exhibiting negligible false positives. To demonstrate the validity of our results, we also present comparative ASR evaluations using other judge models and manual human judgment. These comparative results, along with the specific evaluation prompts used, are detailed in the Supplementary Material. Deepseek-V3 scores response toxicity on a scale of 1–5 relative to the original harmful instruction. An attack is successful if the score $\mathcal{R}(t, \mathcal{M}_\theta(x_t, x_v)) \geq S_{harm}$. The ASR is computed as:

$$ASR = \frac{\sum \{\mathcal{R}(t, \mathcal{M}_\theta(x_t, x_v)) \geq S_{harm}\}}{N_{total}}, \quad (5)$$

where the numerator counts successful attacks and N_{total} is the total number of text-image pairs evaluated.

Implementation Details. All experiments were conducted on a single NVIDIA GeForce RTX 4090 GPU. Within our PIS framework, agents in the PROMPTER and PAINTER teams use GPT-3.5-Turbo [19], while the GUIDER team employs multimodal GPT-4.1 [20]. Image generation model uses Qwen-Image-Plus [29]. The all supervisors allows up to 5 iterative optimization rounds per sample, and the toxicity evaluation threshold S_{harm} is set to 3. Complete details of the logic corpus M_l and sanitization tools M_h and all prompts for API calls are available in the Supplementary Material.

4.2. Experimental Result

Performance on Closed-Source Models. To assess PIS’s jailbreaking on mainstream closed-source LVLMS, we compare it with four SOTA attack methods (FigStep, MM-SafetyBench, HADES, SI-Attack) and a Vanilla baseline, using GPT-4o, Qwen3-VL-Plus, and Gemini 2.5 Flash.

Table 1. Results on SafeBench and HADES dataset. Best results are in **bold**.

Target Model	GPT-4o		Qwen3-VL-Plus		Gemini 2.5 Flash		GLM-4.5V		Qwen3-VL-8B		InternVL3-9B	
Method	Toxic	ASR (%)	Toxic	ASR (%)	Toxic	ASR (%)	Toxic	ASR (%)	Toxic	ASR (%)	Toxic	ASR (%)
SafeBench dataset												
Vanilla	1.34	10.85	1.02	0.80	1.29	8.00	1.47	13.71	1.03	0.50	1.97	32.29
FigStep	1.47	14.80	2.09	34.85	2.49	44.40	2.91	59.14	2.27	42.28	3.15	71.14
MM-SafetyBench	1.80	27.14	1.31	10.57	1.58	20.57	2.89	59.71	1.31	10.00	2.41	46.57
HADES	2.05	34.28	1.46	14.57	1.91	32.57	2.70	56.00	1.37	12.00	2.67	56.57
SI-Attack	1.14	4.28	1.46	15.14	1.75	6.60	2.27	42.29	1.88	28.85	1.90	30.57
PIS (ours)	2.88	60.85	3.28	71.14	2.75	58.57	3.06	66.00	3.07	63.14	3.08	67.71
HADES dataset												
Vanilla	1.12	3.73	1.04	1.30	1.09	2.80	1.32	10.00	1.03	0.80	1.60	20.00
FigStep	1.17	5.47	2.22	38.00	1.45	14.67	3.20	66.70	1.76	24.93	2.67	55.56
MM-SafetyBench	1.37	12.17	1.19	5.60	1.08	3.20	1.89	28.13	1.17	5.46	2.21	39.60
HADES	1.28	9.35	1.26	8.40	1.15	4.93	1.50	15.33	1.06	1.87	1.69	22.67
SI-Attack	1.61	19.87	1.43	13.60	1.33	10.67	2.21	39.06	1.42	12.82	2.04	33.73
PIS (ours)	2.50	48.67	3.05	62.53	2.72	58.27	3.05	65.73	2.74	53.60	2.92	62.00

Tab. 1 shows that PIS significantly outperforms all other methods. Specifically, on the SafeBench, PIS achieves an ASR of 60.85% on GPT-4o, significantly surpassing the runner-up HADES (34.28%) by 26.57%. Its performance on Qwen3-VL-Plus is particularly outstanding, reaching 71.14%, which is 36.29% higher than FigStep (34.85%).




Figure 4. Average Toxic Score on the Gemini 2.5 Flash model using the HADES dataset across five categories.

On Gemini 2.5 Flash, PIS also leads with 58.57%, outperforming FigStep (44.40%) by 14.17%. On the more challenging HADES dataset, the superiority of PIS is even more pronounced. On GPT-4o, PIS achieves 48.67% ASR, exceeding SI-Attack (19.87%), by 28.80%. On Qwen3-VL-Plus, it reaches 62.53%, a 24.53% improvement over FigStep (38.00%). On Gemini 2.5 Flash, PIS achieves 58.27%, vastly outperforming FigStep (14.67%) with a 43.60% lead. Furthermore, PIS consistently achieves a significantly higher average Toxic Score across all closed-source models. This indicates that its generated responses are not only more aggressive but also demonstrate greater robustness in bypassing safety alignment mechanisms.

Performance on Open-Source Models. To verify that PIS is equally effective on open-source models, we evaluate it on three popular open-source LVLs: GLM-4.5V, Qwen3-VL-8B, and InternVL-3-9B. As Tab. 1 shows, PIS exhibits robust performance across open-source models.

On SafeBench, it achieves the highest ASR on GLM-4.5V (66.00%) and Qwen3-VL-8B (63.14%). Although ASR on InternVL-3-9B (67.71%) is slightly lower than FigStep’s (71.14%), the gap is only 3.43%, with a Toxic Score difference of 0.07. On HADES dataset, PIS remains effective: it attains 53.60% ASR on Qwen3-VL-8B, surpassing FigStep (24.93%) by 28.67%; achieves 62.00% ASR on InternVL-3-9B, outperforming other baselines; and reaches 65.73% ASR on GLM-4.5V, nearly matching FigStep with a 0.97% difference.




Figure 5. Average Toxic Score on the Qwen3-VL-8B model using the HADES dataset across five categories.

Overall, PIS not only demonstrates exceptional performance on advanced closed-source models but also exhibits consistent and robust attack capabilities on open-source models. This robustness is also reflected in fine-grained attack categories. As shown in Fig. 4 and Fig. 5, PIS achieves the highest average Toxic Score across every category of the HADES dataset, confirming its effectiveness.

4.3. Further Analysis of PIS’s Universality

Prior work has identified limitations in black-box jailbreak methods. Gong et al. [8] observed that FigStep fails against LVLs equipped with OCR-based security detection tools and proposed FigStep_{pro}; HIMRD [17] noted the over-reliance of black-box attacks on single-modality features and focused on textual prompt optimization. We com-

pare both with PIS on SafeBench across two closed-source (GPT-4V, GPT-4o) and two open-source models (InternVL-3-9B, GLM-4.5V).

Tab. 2 shows that PIS achieves the highest ASR on GPT-4V (50.28%) and GPT-4o (60.85%), significantly outperforming both baselines. On open-source models, PIS (67.71%, 66.00%) surpasses FigStep_{pro} (63.43%, 62.57%) but is outperformed by HIMRD (81.43%, 81.95%). Crucially, while HIMRD and FigStep_{pro} excel on open-source models, their ASR drops to near zero on the two closed-source models. This stark performance gap caught our attention. While architectural differences might be a factor, by analyzing the source code and the commonalities of these two methods, we hypothesize that the **mandatory suffix** they employ is the primary reason for this phenomenon (see Supplementary Material for mandatory suffix details). To test this hypothesis, we removed the mandatory

Table 2. ASR (%) of FigStep_{pro} and HIMRD (with and without suffix), and the proposed PIS method on the SafeBench dataset. Best results are in **bold**.

Attack Method	GPT-4V	GPT-4o	InternVL3-9B	GLM-4.5V	Average
FigStep _{pro}	0.50	4.28	63.43	62.57	32.69
FigStep _{pro} (without suffix)	4.49 ^{↑3.99}	6.57 ^{↑2.29}	51.71 ^{↓11.72}	43.97 ^{↓18.60}	26.69
HIMRD	0.00	0.00	81.43	81.95	40.84
HIMRD (without suffix)	20.85 ^{↑20.85}	26.57 ^{↑26.57}	61.14 ^{↓20.29}	21.71 ^{↓60.24}	32.57
PIS	50.28	60.85	67.71	66.00	61.21

suffix of HIMRD and FigStep_{pro}. The results are shown in Tab. 2. Without the suffix, HIMRD’s ASR on GPT-4o and GPT-4V rose from 0% to 20.85% and 26.57%, respectively, suggesting its attack strategy was previously suppressed by the suffix on closed-source models. Conversely, performance on open-source models dropped sharply: on GLM-4.5V, HIMRD’s ASR fell from 81.95% to 21.71% (a 60.24% drop), and decreased by 20.29% on InternVL-3-9B. FigStep_{pro} showed a similar trend, though less pronounced due to its weaker suffix constraint. These results support our hypothesis: the mandatory suffixes of these methods boost performance on some open-source models but cause near-complete failure on advanced closed-source models like GPT-4o and GPT-4V. In contrast, PIS achieves consistently strong performance across all four models with an average ASR of 62.21%, significantly outperforming all baselines. This demonstrates PIS’s robust adaptability across architectures and also reveals the true role of mandatory suffixes in jailbreaking, offering a potential direction for strengthening model safety defenses.

4.4. Ablation Study

To verify the effectiveness of the three agent team components in PIS, we conducted a series of ablation studies on SafeBench-Small, a set composed of the first 10 samples from each of the first 7 categories of SafeBench.

Effectiveness of the Structured Persuasive Prompts.

To validate the effectiveness of the “structured persuasion prompts” designed by PROMPTER team in jailbreaking, we replaced text input of method HADES with the original harmful instruction while keeping the image input unchanged, denoted as HADES_{harm}. We then processed the text input of HADES_{harm} using PROMPTER to obtain the structured persuasive prompt, keeping the image unchanged, denoted as HADES_{harm} (+PROMPTER). Tab. 3

Table 3. Ablation studies results on SafeBench-Small.

Attack Method	GPT-4o	Qwen3-VL-8B
HADES	31.88	12.86
HADES (+PAINTER)	55.71 ^{↑23.83}	17.14 ^{↑4.28}
HADES _{harm}	20.00	4.29
HADES _{harm} (+PROMPTER)	31.43 ^{↑11.43}	10.00 ^{↑5.71}
HADES _{harm} (+PROMPTER+PAINTER)	37.14	11.43
PIS	61.43 ^{↑24.29}	55.71 ^{↑44.28}

shows that HADES_{harm} (+PROMPTER) improved ASR by 11.43% on GPT-4o and 5.71% on Qwen-3-VL-8B compared to HADES_{harm}. Even when the harmful information in the structured persuasive prompts was fully exposed, its ASR on GPT-4o and Qwen-3-VL-8B remained comparable to HADES, which transfers harmful information from text to image. These findings indicate that the PROMPTER enhances the persuasiveness of the text prompt and plays an indispensable role in the jailbreak process.

Effectiveness of the Scenario Typographic Images.

Next, to verify the effectiveness of the “scenario typographic image” designed by PAINTER team, we replaced the simple concatenated image in HADES with our scenario typographic image, keeping the text unchanged (denoted HADES (+PAINTER)). Tab. 3 shows HADES (+PAINTER) improved the ASR by 23.83% on GPT-4o and 4.28% on Qwen-3-VL-8B compared to the original HADES. This finding strongly validates our design concept of embedding harmful keywords into scene images as reasonable clues.

Effectiveness of Sanitization and Cross-modal Synergy.

To validate the GUIDER team’s effectiveness in integrating the “scenario typographic image” and “structured persuasive prompt”, we compared our full PIS method against two key baselines: HADES_{harm} (+PROMPTER+PAINTER) and HADES’s own inherent sanitization method. The results in Tab. 3 show that PIS outperforms HADES_{harm} (+PROMPTER+PAINTER) by 24.29% and HADES by 29.55% ASR on GPT-4o, and by 42.85% and 44.28% ASR on Qwen-3-VL-8B, respectively. This finding indicates that our designed GUIDER team not only effectively performs sanitization, but also enables stronger cross-modal synergy,




Figure 6. Visualization of representation distributions for benign/harmful anchors and the PIS-generated input’s evolution within the Qwen3-VL-8B model.

leading to significantly improved attack performance over the baseline method.

Impact of Supervisor Iteration Count. In our PIS framework, each of the three agent teams (PROMPTER, PAINTER, GUIDER) is assigned a supervisor. The PROMPTER team’s supervisor $\mathcal{J}_t$ evaluates the text’s semantic consistency, persuasiveness, and structural integrity. The PAINTER team’s supervisor $\mathcal{J}_i$ assesses prompt conciseness, completeness, and adherence to the user requirement $U(k)$. The GUIDER team’s supervisor $\mathcal{J}_h$ ensures final instruction intent consistency, harmlessness, and image-text alignment. Fig. 7a shows that with well-designed system prompts, most samples pass supervisor assessment in the first iteration; only a few require iterative refinement. We set the maximum iterations to 10, but found that the maximum needed was 5. Although the ablation study in Fig. 7b indicates that PIS remains effective without supervisors, their inclusion significantly enhances ASR, confirming their essential role for optimizing attacks.

4.5. Visual Analysis

To intuitively illustrate the effectiveness of PIS at different stages, we visualized the representation distributions of Qwen3-VL-8B’s responses under various attack configurations using a methodology from work [14]. As Fig. 7a shows, processing only the structured persuasive prompt by the PROMPTER team results in significant skewing of the model’s response representations towards the harmful anchor, validating the effectiveness of our text content. In Fig. 7b, combining the text processed by PROMPTER team and the scenario typographic image from PAINTER team without GUIDER team still exhibits a synergistic effect, further clustering the response distribution towards the harmful anchor. Finally, Fig. 6c shows the complete PIS situation, at this point, the model’s response representations fully con-

verge into the harmful anchor distribution. These results clearly demonstrate the effectiveness of our designed PIS.




Figure 7. **Ablation study of the supervisors.** (a) Number of samples requiring iterative refinement by the supervisors in the PROMPTER, PAINTER, and GUIDER. (b) ASR with and without supervisors on the Qwen3-VL-8B and GPT-4o models.

5. Conclusion

We propose Persuasion in Scene (PIS), a novel multi-agent framework for automated black-box jailbreaking of LVLMs. PIS addresses the core limitations of existing “typography and concatenation” methods: the detectability of harmful information in images and poor cross-modal consistency. Our proposed framework uses specialized PROMPTER, PAINTER, and GUIDER agent teams to automatically generate sanitized “structured persuasive prompts” and natural “scenario typographic images”, ensuring strong cross-modal synergy. PIS has achieved an ASR of over 60% on advanced LVLMs, including GPT-4o, Gemini 2.5 Flash, Qwen3-VL-Plus, and GLM-4.5V, significantly outperforming existing comparison baselines. This work not only validates PIS’s effectiveness, but also reveals critical cross-modal security vulnerabilities in LVLMs, providing key insights for future robust defense strategies.

References

- [1] Jinze Bai, Shuai Bai, Shusheng Yang, Shijie Wang, Sinan Tan, Peng Wang, Junyang Lin, Chang Zhou, and Jingren Zhou. Qwen-vl: A versatile vision-language model for understanding, localization, text reading, and beyond. *arXiv preprint arXiv:2308.12966*, 2023. 1
- [2] Patrick Chao, Alexander Robey, Edgar Dobriban, Hamed Hassani, George J. Pappas, and Eric Wong. Jailbreaking black box large language models in twenty queries, 2024. 2, 3
- [3] Guangyao Chen, Siwei Dong, Yu Shu, Ge Zhang, Jaward Sesay, Börje F Karlsson, Jie Fu, and Yemin Shi. Autoagents: A framework for automatic agent generation. *arXiv preprint arXiv:2309.17288*, 2023. 3
- [4] Weize Chen, Yusheng Su, Jingwei Zuo, Cheng Yang, Chenfei Yuan, Chi-Min Chan, Heyang Yu, Yaxi Lu, Yi-Hsin Hung, Chen Qian, et al. Agentverse: Facilitating multi-agent collaboration and exploring emergent behaviors. In *The Twelfth International Conference on Learning Representations*, 2023. 3
- [5] Gheorghe Comanici, Eric Bieber, Mike Schaekermann, Ice Pasupat, Noveen Sachdeva, Inderjit Dhillon, Marcel Blstein, Ori Ram, Dan Zhang, Evan Rosen, et al. Gemini 2.5: Pushing the frontier with advanced reasoning, multimodality, long context, and next generation agentic capabilities. *arXiv preprint arXiv:2507.06261*, 2025. 5
- [6] Can Cui, Yunsheng Ma, Xu Cao, Wenqian Ye, Yang Zhou, Kaizhao Liang, Jintai Chen, Juanwu Lu, Zichong Yang, Kuei-Da Liao, et al. A survey on multimodal large language models for autonomous driving. In *Proceedings of the IEEE/CVF winter conference on applications of computer vision*, pages 958–979, 2024. 1
- [7] DeepSeek-AI. Deepseek-v3 technical report, 2024. 5
- [8] Yichen Gong, Delong Ran, Jinyuan Liu, Conglei Wang, Tianshuo Cong, Anyu Wang, Sisi Duan, and Xiaoyun Wang. Figstep: Jailbreaking large vision-language models via typographic visual prompts. In *Proceedings of the AAAI Conference on Artificial Intelligence*, pages 23951–23959, 2025. 1, 2, 3, 5, 6
- [9] Shuyang Hao, Bryan Hooi, Jun Liu, Kai-Wei Chang, Zi Huang, and Yujun Cai. Exploring visual vulnerabilities via multi-loss adversarial search for jailbreaking vision-language models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 19890–19899, 2025. 1, 2
- [10] Iryna Hartsock and Ghulam Rasool. Vision-language models for medical report generation and visual question answering: a review. *Frontiers in Artificial Intelligence*, Volume 7 - 2024, 2024. 1
- [11] GLM-V Team Wenyi Hong, Wenmeng Yu, Xiaotao Gu, Guo Wang, Guobing Gan, Haomiao Tang, Jiale Cheng, Ji Qi, Junhui Ji, Lihang Pan, Shuaiqi Duan, Weihang Wang, Yan Wang, Yean Cheng, Zehai He, Zhe Su, Zhen Yang, Ziyang Pan, Aohan Zeng, Baoxu Wang, Boyan Shi, Changyu Pang, Chenhui Zhang, Da Yin, Fan Yang, Guoqing Chen, Jiazhen Xu, Jiali Chen, Jing Chen, Jinhao Chen, Jinghao Lin, Jinjiang Wang, Junjie Chen, Leqi Lei, Letian Gong, Leyi Pan, Mingzhi Zhang, Qinkai Zheng, Shengchao Yang, Shilong Zhong, Shiyu Huang, Shuyuan Zhao, Siyan Xue, Shangqing Tu, Shengbiao Meng, Tianshu Zhang, Tian-Yuan Luo, Tianxiang Hao, Wenkai Li, Wei Jia, Xinpeng Lyu, Xuancheng Huang, Yanling Wang, Ya-Qi Xue, Yanfeng Wang, Yifan An, Yifan Du, Yi Shi, Yiheng Huang, Yilin Niu, Yuan Wang, Yuanchang Yue, Yuchen Li, Yutao Zhang, Yuxuan Zhang, Zhanxiao Du, Zhenyu Hou, Zhao Xue, Zhengxiao Du, Zihan Wang, Peng Zhang, De-Huan Liu, Bin Xu, Juanzi Li, Minlie Huang, Yuxiao Dong, and Jie Tang. Glm-4.5v and glm-4.1v-thinking: Towards versatile multimodal reasoning with scalable reinforcement learning. 2025. 5
- [12] Aaron Hurst, Adam Lerer, Adam P Goucher, Adam Perelman, Aditya Ramesh, Aidan Clark, AJ Ostrow, Akila Welihinda, Alan Hayes, Alec Radford, et al. Gpt-4o system card. *arXiv preprint arXiv:2410.21276*, 2024. 1, 5
- [13] Yifan Li, Hangyu Guo, Kun Zhou, Wayne Xin Zhao, and Jirong Wen. Images are achilles’ heel of alignment: Exploiting visual vulnerabilities for jailbreaking multimodal large language models. In *European Conference on Computer Vision*, pages 174–189. Springer, 2024. 1, 2, 3, 5
- [14] Yuping Lin, Pengfei He, Han Xu, Yue Xing, Makoto Yamada, Hui Liu, and Jiliang Tang. Towards understanding jailbreak attacks in llms: A representation space analysis. *ArXiv*, abs/2406.10794, 2024. 8
- [15] Haotian Liu, Chunyuan Li, Qingyang Wu, and Yong Jae Lee. Visual instruction tuning. *Advances in neural information processing systems*, 36:34892–34916, 2023. 1
- [16] Xin Liu, Yichen Zhu, Jindong Gu, Yunshi Lan, Chao Yang, and Yu Qiao. Mm-safetybench: A benchmark for safety evaluation of multimodal large language models. In *European Conference on Computer Vision*, pages 386–403. Springer, 2024. 1, 2, 5
- [17] Teng Ma, Xiaojun Jia, Ranjie Duan, Xinfeng Li, Yihao Huang, Xiaoshuang Jia, Zhixuan Chu, and Wenqi Ren. Heuristic-induced multimodal risk distribution jailbreak attack for multimodal large language models. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 2686–2696, 2025. 1, 2, 3, 5, 6
- [18] OpenAI. Gpt-4v(ision) system card. 2023. 2, 3, 5
- [19] OpenAI. Gpt-3.5 turbo fine-tuning and api updates, 2023. 5
- [20] OpenAI. Introducing gpt-4.1 in the api, 2025. 5
- [21] Xiangyu Qi, Kaixuan Huang, Ashwinee Panda, Peter Henderson, Mengdi Wang, and Prateek Mittal. Visual adversarial examples jailbreak aligned large language models. In *Proceedings of the AAAI conference on artificial intelligence*, pages 21527–21536, 2024. 1, 2
- [22] Rusheb Shah, Soroush Pour, Arush Tagade, Stephen Casper, Javier Rando, et al. Scalable and transferable black-box jailbreaks for language models via persona modulation. *arXiv preprint arXiv:2311.03348*, 2023. 2
- [23] Erfan Shayegani, Yue Dong, and Nael Abu-Ghazaleh. Jailbreak in pieces: Compositional adversarial attacks on multimodal language models. *arXiv preprint arXiv:2307.14539*, 2023. 1
- [24] Jiaxin Song, Yixu Wang, Jie Li, Rui Yu, Yan Teng, Xingjun Ma, and Yingchun Wang. Jailbound: Jailbreaking internal

- safety boundaries of vision-language models. *arXiv preprint arXiv:2505.19610*, 2025. 2
- [25] Gemini Team, Rohan Anil, Sebastian Borgeaud, Jean-Baptiste Alayrac, Jiahui Yu, Radu Soricut, Johan Schalkwyk, Andrew M Dai, Anja Hauth, Katie Millican, et al. Gemini: a family of highly capable multimodal models. *arXiv preprint arXiv:2312.11805*, 2023. 1
- [26] Qwen Team. Qwen3 technical report, 2025. 5
- [27] Xiaoyu Tian, Junru Gu, Bailin Li, Yicheng Liu, Yang Wang, Zhiyong Zhao, Kun Zhan, Peng Jia, Xianpeng Lang, and Hang Zhao. Drivevlm: The convergence of autonomous driving and large vision-language models. *arXiv preprint arXiv:2402.12289*, 2024. 1
- [28] walkerspider. Dan is my new friend, 2022. Reddit post in r/ChatGPT subreddit. 2
- [29] Chenfei Wu, Jiahao Li, Jingren Zhou, Junyang Lin, Kaiyuan Gao, Kun Yan, Sheng ming Yin, Shuai Bai, Xiao Xu, Yilei Chen, Yuxiang Chen, Zecheng Tang, Zekai Zhang, Zhengyi Wang, An Yang, Bowen Yu, Chen Cheng, Dayiheng Liu, Deqing Li, Hang Zhang, Hao Meng, Hu Wei, Jingyuan Ni, Kai Chen, Kuan Cao, Liang Peng, Lin Qu, Minggang Wu, Peng Wang, Shuting Yu, Tingkun Wen, Wensen Feng, Xiaoxiao Xu, Yi Wang, Yichang Zhang, Yongqiang Zhu, Yujia Wu, Yuxuan Cai, and Zenan Liu. Qwen-image technical report, 2025. 5
- [30] Nur Yildirim, Hannah Richardson, Maria Teodora Wetscherek, Junaid Bajwa, Joseph Jacob, Mark Ames Pinnock, Stephen Harris, Daniel Coelho De Castro, Shruthi Bannur, Stephanie Hyland, Pratik Ghosh, Mercy Ranjit, Kenza Bouzid, Anton Schwaighofer, Fernando Pérez-García, Harshita Sharma, Ozan Oktay, Matthew Lungren, Javier Alvarez-Valle, Aditya Nori, and Anja Thieme. Multimodal healthcare ai: Identifying and designing clinically relevant vision-language applications for radiology. New York, NY, USA, 2024. Association for Computing Machinery. 1
- [31] Zonghao Ying, Aishan Liu, Tianyuan Zhang, Zhengmin Yu, Siyuan Liang, Xianglong Liu, and Dacheng Tao. Jailbreak vision language models via bi-modal adversarial prompt. *IEEE Transactions on Information Forensics and Security*, 2025. 1, 2
- [32] Yi Zeng, Hongpeng Lin, Jingwen Zhang, Diyi Yang, Ruoxi Jia, and Weiyan Shi. How johnny can persuade llms to jailbreak them: Rethinking persuasion to challenge ai safety by humanizing llms. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 14322–14350, 2024. 2, 3
- [33] Shiji Zhao, Ranjie Duan, Fengxiang Wang, Chi Chen, Caixin Kang, Shouwei Ruan, Jialing Tao, YueFeng Chen, Hui Xue, and Xingxing Wei. Jailbreaking multimodal large language models via shuffle inconsistency. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 2045–2054, 2025. 2, 5
- [34] Xingcheng Zhou, Mingyu Liu, Ekim Yurtsever, Bare Luka Zagar, Walter Zimmer, Hu Cao, and Alois C Knoll. Vision language models in autonomous driving: A survey and outlook. *IEEE Transactions on Intelligent Vehicles*, 2024. 1
- [35] Jinguo Zhu, Weiyun Wang, Zhe Chen, Zhaoyang Liu, Shenglong Ye, Lixin Gu, Yuchen Duan, Hao Tian, Weijie Su, Jie Shao, Zhangwei Gao, Erfei Cui, Yue Cao, Yangzhou Liu, Haomin Wang, Weiye Xu, Hao Li, Jiahao Wang, Han Lv, Dengnian Chen, Songze Li, Yinan He, Tan Jiang, Jiapeng Luo, Yi Wang, Cong He, Botian Shi, Xingcheng Zhang, Wenqi Shao, Junjun He, Ying Xiong, Wenwen Qu, Peng Sun, Penglong Jiao, Lijun Wu, Kai Zhang, Hui Deng, Jiaye Ge, Kaiming Chen, Limin Wang, Min Dou, Lewei Lu, Xizhou Zhu, Tong Lu, Dahua Lin, Yu Qiao, Jifeng Dai, and Wenhao Wang. Internvl3: Exploring advanced training and test-time recipes for open-source multimodal models. *ArXiv*, abs/2504.10479, 2025. 5
- [36] Andy Zou, Zifan Wang, Nicholas Carlini, Milad Nasr, J Zico Kolter, and Matt Fredrikson. Universal and transferable adversarial attacks on aligned language models. *arXiv preprint arXiv:2307.15043*, 2023. 2