Unified integration of pathology foundation models for scalable histopathology analysis

Wenhui Lei[1‡], Yusheng Tan[2‡], Anqi Li[2‡], Hengrui Tian[1], Ruiying Li[3], Zhenqun Jiang[1], Fang Yan[4], Hanyu Chen[*5], Xiaofan Zhang[*1,6], Shaoting Zhang[*1,7]

[1]*Shanghai Jiao Tong University, Shanghai, China*

[2]*Washington University in St. Louis, St. Louis, USA*

[3]*University of Science and Technology Beijing, Beijing, China*

[4]*Shanghai Artificial Intelligence Laboratory, Shanghai, China*

[5]*Department of Surgical Oncology and General Surgery, Key Laboratory of Precision Diagnosis and Treatment of Gastrointestinal Tumours, Ministry of Education, The First Hospital of China Medical University, Liaoning, China*

[6]*Shanghai Innovation Institute, Shanghai, China*

[7]*Sensetime Research, Shanghai, China*

‡ *Contributed Equally*

***Corresponding author**: Hanyu Chen (hychen88@cmu.edu.cn), Xiaofan Zhang (xiaofan.zhang@sjtu.edu.cn), Shaoting Zhang (shaoting.zhang@sjtu.edu.cn)*

arXiv:2503.00736v4 [cs.CV] 13 Feb 2026

Foundation models have substantially advanced computational pathology by learning transferable visual representations from large histological datasets, yet their performance varies widely across tasks due to differences in training data composition and reliance on proprietary datasets that cannot be cumulatively expanded. Existing efforts to combine foundation models through offline distillation partially mitigate this issue but require dedicated distillation data and repeated retraining to integrate new models. Here we present Shazam, an online integration model that adaptively combines multiple pretrained pathology foundation models within a unified and scalable representation learning paradigm. Our findings show that fusing multi-level features through adaptive expert weighting and online distillation enables efficient consolidation of complementary model strengths without additional pretraining. Across spatial transcriptomics prediction, survival prognosis, tile-level classification, and visual question answering, Shazam consistently outperforms strong individual models, demonstrating that online model integration provides a practical and extensible strategy for advancing computational pathology.

Introduction

Histopathology remains the cornerstone of cancer diagnosis and treatment planning, providing direct visualization of tissue architecture, cellular morphology, and tumor–microenvironment interactions that guide clinical decision-making. Advances in whole-slide imaging (WSI) have transformed this traditionally microscope-based discipline into a scalable digital workflow, enabling the systematic capture, storage, and analysis of gigapixel-resolution slides. This digital transition has laid the foundation for modern computational pathology (CPath), where artificial intelligence systems can interrogate WSI data to support diagnostic, predictive, and prognostic tasks[1,2]. Building on this momentum, recent pathology foundation models have introduced a new paradigm for CPath. By pretraining on millions of tiles extracted from large cohorts of WSIs across diverse organs and cancer types, models such as Virchow[3], UNI[4], and H-Optimus-1[5] learn broad and transferable representations that enhance data efficiency and improve downstream performance across a wide spectrum of clinical tasks.

Despite these advances, recent benchmarking studies have revealed substantial performance variability among pathology foundation models across tasks[6–8]. Much of this inconsistency reflects differences in the scale, diversity, and institutional origins of the pretraining data, and no single model consistently performs well across diagnostic, prognostic, molecular, and region-level tasks. Compounding this challenge, many high-performing models rely on large proprietary datasets that cannot be shared due to privacy and institutional constraints[4,5,9]. This lack of data accessibility prevents the research community from steadily expanding the pretraining corpus over time, constraining the growth in data diversity that has supported advances in natural image and language foundation models[10–12].

One line of recent work has sought to mitigate limited access to large-scale pathology training data by combining existing foundation models through offline knowledge distillation, exemplified by GPFM[13]. In this framework, multiple pretrained teacher models first generate features on a curated distillation dataset, and a student model is trained to learn from these signals through expert distillation, complemented by self-distillation that aligns local and global representations. While this strategy provides a practical means of integrating knowledge from several models, it faces several inherent limitations. First, because the distillation process is performed offline, incorporating newly developed foundation models requires retraining the student model from scratch, making rapid integration of emerging advances difficult. Second, the effectiveness of offline distillation depends heavily on the scale and diversity of the available distillation dataset, which in practice is restricted to publicly accessible or newly assembled collections. This constraint limits the range of morphologic and clinical variation that the student model can encounter during distillation, reducing its ability to fully capture the complementary behaviors of diverse teacher models[14]. Third, offline distillation provides

limited flexibility for task-specific adaptation, since the student model learns from teacher outputs generated in a task-agnostic manner and cannot selectively emphasize teacher knowledge that is particularly relevant for a specific downstream prediction target. These constraints highlight the need for an alternative strategy that can directly leverage multiple foundation models in a more flexible and adaptive way, without relying on repeated offline retraining or large-scale distillation corpora.

To directly address these challenges, we developed Shazam, a unified ensemble model that integrates multiple pathology foundation models within a single flexible inference system (Fig. 1a). In contrast to offline distillation approaches that require assembling a dedicated distillation dataset and retraining a new model, Shazam performs task-specific online integration. It combines multi-level representations from diverse pathology foundation models through a lightweight Mixture-of-Experts (MoE)[15] mechanism and adapts the fusion process directly to the target task. This design enables Shazam to extract the most informative elements from each model without the need to restart pretraining or construct large-scale distillation corpora whenever new foundation models become available. We evaluate Shazam across a wide range of downstream applications, including spatial transcriptomics prediction, whole-slide survival prognosis, tile-level classification, and visual question answering. These tasks span molecular, regional, and slide-level endpoints and together reflect the breadth of clinical challenges encountered in computational pathology. Across these evaluations, Shazam demonstrates consistent and robust performance. In total, it achieves an average ranking score of 1.17 across 30 benchmarks, substantially outperforming the second-best model, Virchow2[16], which attains an average of 3.20. By enabling task-adaptive, online integration of heterogeneous pretrained models, Shazam provides a practical mechanism for consolidating complementary representations without requiring additional large-scale pretraining or privileged data access. More broadly, this paradigm suggests a shift in pathology foundation modeling from isolated, monolithic models toward modular and compositional systems, offering a scalable path for continual methodological evolution under realistic clinical and institutional constraints.

Although Shazam led performance on most benchmarks, its advantage narrowed in datasets where a single baseline model exhibited a pronounced task-specific strength. For instance, on UniToPatho, Virchow2 achieved a F1 score of 0.820, noticeably higher than the other individual models, and Shazam obtained 0.803, reflecting that the ensemble can inherit attenuated signals when several constituent models underperform. Even in such settings, Shazam remained consistently competitive across metrics, maintaining reliable performance despite variations in stain, scanner, and tissue preparation. This stability highlights Shazam's capacity to generalize fine-grained morphological patterns at the tile level, where subtle artifacts, microenvironmental cues, and local histologic variability pose the greatest challenges for computational pathology models.

4 Shazam improves pathology VQA

Visual Question Answering (VQA) aims to answer natural-language queries based on visual evidence and has emerged as a promising framework for assisting clinicians in retrieving task-specific diagnostic information from histopathology images. To examine Shazam's capability in this setting, we used the PathVQA[39] dataset, which is among the largest resources for pathology VQA. The dataset contains 32,799 image–question–answer triplets, divided into a training set of 16,400 samples, a validation set of 9,840 samples, and a test set of 6,560 samples. Overall performance across the evaluated pathology foundation models is presented in Fig. 4c, with detailed results provided in Supplementary Table 19.

Shazam achieved the highest accuracy of 0.575, exceeding the performance of the second-best model, UNI 2, by 1%. This result is notable because the foundation models underlying Shazam were trained exclusively with image-based self-supervision and therefore received no explicit language or cross-modal guidance. Despite this limitation, the unified multi-model, multi-scale architecture of Shazam enabled strong performance on a task requiring alignment between visual patterns and linguistic queries. These findings indicate that Shazam's integrative feature representation improves not only the discriminability of image-based features but also their transferability to cross-modal reasoning tasks. As a result, Shazam is able to extract visual cues from histopathology images that are relevant to the semantics of each question.

We further provide qualitative comparisons in Fig. 4d, where we illustrate a representative query image, question, and predicted answers across different foundation models. Shazam and UNI 2 consistently produced more reliable and clinically meaningful responses than the other evaluated models, supporting the quantitative advantages observed in the benchmark analysis.

3.7% decrease). Together, these results support our hypothesis that prioritizing task-aligned foundation models yields more effective supervision than treating all teachers equally, and that integrating such models within Shazam leads to measurable and robust performance improvements.

We next quantified the contribution of feature representations at different levels of semantic abstraction for whole-slide survival prediction on TCGA-COADREAD (Fig. 5b; Supplementary Table 21). Using features from a single level yielded C-index values in the range of 0.680–0.689. Pairwise combinations improved performance by up to 0.022 in absolute C-index, with the low- and high-level combination achieving 0.702. Aggregating low-, mid-, and high-level features achieved the highest C-index of 0.710, representing an absolute improvement of 0.030 over the best single-level configuration and 0.008 over the best pairwise combination. These results demonstrate that representations encoding local morphology, intermediate tissue organization, and global contextual patterns capture complementary prognostic information.

Finally, we evaluated the effect of the Mixture-of-Experts (MoE) module on CRC-100K and CCRCC-TCGA-HEL classification tasks (Fig. 5c; Supplementary Tables 22-23). On CRC-100K, incorporating MoE improved balanced accuracy from 0.955 to 0.960 (+0.005), weighted F1 from 0.951 to 0.956 (+0.005), and top-1 accuracy from 0.970 to 0.972 (+0.002). On CCRCC-TCGA-HEL, the corresponding improvements were smaller but consistent, with weighted F1 increasing from 0.973 to 0.975 (+0.002) and top-1 accuracy from 0.973 to 0.975 (+0.002). These results suggest that adaptive expert routing provides measurable but dataset-dependent gains over uniform feature aggregation.

Together, these ablation studies indicate that Shazam's performance gains arise from the interaction of strong but uneven teacher supervision, complementary semantic feature representations, and adaptive aggregation. While the largest gains stem from retaining informative teachers and multi-level features, MoE further refines performance by enabling task-adaptive fusion of heterogeneous representations.

Discussion

In this work, we introduce Shazam, a unified pathology foundation model that consolidates the capabilities of multiple pretrained expert models without requiring access to their training data or performing any form of large-scale distillation pretraining. Previous efforts such as GPFM attempted to merge knowledge through offline distillation, yet their effectiveness is fundamentally limited by the diversity and scale of the available distillation corpus. Shazam adopts a fundamentally different integration strategy that operates in an online and task-specific manner, combining pretrained pathology foundation models through multi-level feature fusion and an adaptive expert-weighting mechanism. This design allows Shazam to continuously benefit from advances in pathology foundation models and to flexibly incorporate the most recent and most powerful models for any specific downstream task. With this ability, Shazam demonstrates consistently strong performance

across four major categories of computational pathology, including spatial transcriptomics prediction, survival analysis, tile-level classification, and visual question answering in pathology. Across a total of 30 benchmark tasks, Shazam surpasses leading individual models such as UNI 2 and Virchow2 in the majority of evaluations.

The rapid development of pathology foundation models in recent years has markedly advanced the field, yet their pretraining pipelines continue to rely heavily on large, privately curated datasets whose composition varies substantially across institutions. These differences lead individual models to specialize in particular tissue types, staining styles or diagnostic tasks, which makes it difficult to ensure broad generalization across the diverse landscape of computational pathology. Moreover, because the underlying training data of these models are inaccessible, progress in foundation model development cannot accumulate at the data level, and each new model is effectively trained in isolation from its predecessors. Shazam addresses this structural limitation by allowing downstream tasks to selectively integrate the strengths of multiple pretrained models, without requiring access to any of their training corpora. Through this design, Shazam demonstrates robust generalization across heterogeneous tasks, suggesting that online integration of multiple foundation models represents a promising direction for the next generation of computational pathology systems. This is particularly valuable in settings where different models capture distinct aspects of morphological or molecular variation, and where a single architecture is unlikely to dominate across all task domains.

Despite its strong empirical performance, Shazam has several limitations that warrant further investigation. Because Shazam requires all constituent foundation models to perform feature extraction during both training and inference, its computational cost grows approximately linearly with the number of integrated models. This requirement may limit its practicality in settings where low latency or constrained hardware resources are essential. Recent progress in developing lightweight pathology foundation models with substantially fewer parameters yet competitive performance provides a promising avenue for reducing this overhead[40], and future versions of Shazam could incorporate such models to achieve a more favorable efficiency–performance balance. Another limitation lies in the breadth of external validation. Although Shazam demonstrates consistent gains across a wide range of internal benchmarks, the evaluation on independent datasets remains limited. Expanding external validation to additional cohorts and institutions will be important for fully establishing its robustness under real-world distributional variability. Continued exploration along these directions may further enhance the scalability, generalizability, and clinical readiness of Shazam.

Online Methods

Framework

the 5 folds are presented in Supplementary Table 6.

Glioma Survival Prediction.

We merge TCGA-GBM (370 cases, 823 WSIs) and TCGA-LGG[47] (460 cases, 778 WSIs) into a combined glioma dataset of 830 cases (1,601 WSIs). Per fold, 667 cases are used for training and 163 for testing. The results across the 5 folds are presented in Supplementary Table 6.

Kidney Renal Clear Cell Carcinoma Survival Prediction.

The TCGA-KIRC dataset[47] comprises 498 cases (504 WSIs). Following the same partitioning strategy, each fold assigns 401 cases to training and 97 cases to testing. The results across the 5 folds are presented in Supplementary Table 6.

Lung Adenocarcinoma Survival Prediction.

For TCGA-LUAD[47], we use 455 cases (518 WSIs), splitting each fold into 366 training cases and 89 testing cases. The results across the 5 folds are presented in Supplementary Table 7.

Lung Squamous Cell Carcinoma Survival Prediction.

The TCGA-LUSC dataset[47] comprises 452 cases (484 WSIs). Using the same stratified splitting approach, 365 cases are assigned to training and 87 to testing in each fold. The results across the 5 folds are presented in Supplementary Table 7.

Skin Cutaneous Melanoma Survival Prediction.

The TCGA-SKCM dataset[47] contains 415 cases (456 WSIs). Each fold uses 337 cases for training and 78 for testing. The results across the 5 folds are presented in Supplementary Table 7.

Supplementary Table 1: **Cancer-specific gene lists for HEST-1k (Part 1/2).** Gene prioritization was informed by the FDA-recognized content in OncoKB [44].

Cancer	Gene list
Breast	ALK;APC;AR;ARAF;ARID1A;ATM;BRAF;BRCA1;BRCA2;CCND1;CD274;CDH1;CDK4;CDK6;CDKN2A;CHEK2;CLDN18;CTNNB1; DDR2;DNMT3A;EGFR;EPCAM;ERBB2;ERBB3;ESR1;FBXW7;FGFR1;FGFR2;FGFR3;FGFR4;GNA11;GNAQ;GNAS;HRAS;IDH1;IDH2; JAK1;JAK2;KDR;KEAP1;KIT;KRAS;MAP2K1;MAP2K2;MDM2;MET;MLH1;MSH2;MSH3;MSH6;MTOR;MYC;NF1;NRAS;NRG1;NTRK1; NTRK2;NTRK3;PALB2;PBRM1;PDGFRA;PIK3CA;PIK3R1;PMS2;POLD1;POLE;PTEN;RAF1;RB1;RET;RNF43;ROS1;RSPO2;RSPO3;SDHB; SMAD4;SMARCA4;STK11;TERT;TP53;TSC1;TSC2;VHL;
Bowel	ALK;APC;ARID1A;ATM;BRAF;BRCA1;BRCA2;CCND1;CD274;CDH1;CDK4;CDK6;CDKN2A;CHEK2;CTNNB1;DDR2;DNMT3A;EGFR; EPCAM;ERBB2;ERBB3;FBXW7;FGFR1;FGFR2;FGFR3;FGFR4;GNAS;HRAS;IDH1;IDH2;JAK1;JAK2;KDR;KEAP1;KIT;KRAS;MAP2K1; MDM2;MET;MLH1;MSH2;MSH3;MSH6;MTOR;MYC;NF1;NRAS;NRG1;NTRK1;NTRK2;NTRK3;PALB2;PBRM1;PDGFRA;PIK3CA; PIK3R1;PMS2;POLD1;POLE;PTEN;RAF1;RB1;RET;RNF43;ROS1;RSPO2;RSPO3;SDHB;SMAD4;SMARCA4;STK11;TERT;TP53;
Kidney	CDKN2A;FGFR2;MED12;PIK3CA;TERT;ALK;CDKN2B;FGFR3;MET;PMS2;TFE3;APC;CREBBP;FH;MLH1;POLE;TFEB;AR;CYP2C8; FLCN;MSH2;PTEN;TP53;ARID1A;DHFR;FLT3;MSH6;RAF1;TSC1;ARID2;DPYD;GNAQ;MTHFR;RB1;TSC2;ATIC;DYNC2H1;GNAS; MTOR;RET;UGT1A1;ATRX;EGFR;GSTP1;MYC;ROS1;VEGFA;BAP1;EMSY;HOXB13;NF1;SDHA;VHL;BMPR1A;EPCAM;HRAS;NFE2L2; SDHB;XPC;BRAF;ERBB2;KDM5C;NOTCH2;SDHC;XRCC1;CBR3;ERCC1;KIT;NRAS;SDHD;CCND1;FGF19;KRAS;NTRK1;SETD2;CDA; FGF3;MAP3K1;NTRK2;SMAD4;CDK4;FGF4;MDM2;NTRK3;SMARCA4;CDK6;FGFR1;MDM4;PBRM1;TCEB1;
Brain	ABCC2;ABCC4;ABCG2;ABL1;ABRAXAS1;ALK;APC;AR;ARID1B;ASXL1;ATM;ATR;ATRX;AURKA;BAP1;BARD1;BCL11A;BCL11B; BCORL1;BLM;BRAF;BRCA1;BRCA2;BRD3;BRIP1;C8orf34;CASP3;CASP9;CCDC6;CDH1;CDK1;CDK12;CDK7;CDK8;CDK9;CDKN1C; CDKN2A;CDKN2B;CHEK1;CHEK2;CREBBP;CTNNA1;CUL3;CXCL8;CXCR4;CYP19A1;CYP1B1;CYP2B6;CYP2C8;DDIT3;DDR1;DHFR; DPYD;EGFR;EMSY;EP300;EPCAM;EPHA5;EPHA6;EPHA7;ERBB2;ERCC1;ERCC2;ERCC3;ERCC4;ESR1;ESR2;EWSR1;EZH2;FANCA; FANCC;FANCD2;FANCE;FANCF;FANCG;FANCI;FANCL;FANCM;FGFR1;FGFR2;FGFR3;FLT3;GATA1;GATA4;GEN1;GGH;GSTP1;H1-2; H1-5;H3-3A;H3C2;H3C3;HDAC2;HRAS;IDH1;IDH2;IGF2R;KDR;KIT;KMT2A;KMT2B;KMT2C;KMT2D;KRAS;MAP2K4;MAP3K1; MAP4K1;MAP4K4;MAPK14;MET;MGMT;MLH1;MLH3;MMP14;MMP2;MMP9;MRE11;MSH2;MSH6;MTHFR;MTR;MUTYH;MYCN;NBN; NCOR1;NCOR2;NF1;NOTCH3;NOTCH4;NQO1;NRAS;NTRK1;NTRK2;NTRK3;PALB2;PARP1;PARP2;PARP3;PARP4;PDGFRA;PDGFRB; PIK3C3;PIK3CA;PIK3CD;PML;PMS1;PMS2;POLD1;POLD2;POLE;POLE2;PPM1D;PPP2R2A;PRKDC;PTEN;RAD50;RAD51;RAD51B; RAD51C;RAD51D;RAD52;RAD54B;RAD54L;RB1;RECQL4;RET;ROS1;RPA1;RRM1;SEMA3C;SLC19A1;SLC22A2;SLC28A1;SLCO1B1; SLX4;SMARCB1;SOD2;STAT3;STAT4;STK11;TERT;TGFBR1;TGFBR2;TNFRSF14;TNFRSF17;TP53;TPMT;TSC1;TSC2;UGT1A1;UMPS; WRN;XPC;XRCC1;XRCC2;YAP1;ZBTB16;ZFTA;ZNF217;

Supplementary Table 2: **Cancer-specific gene lists for HEST-1k (Part 2/2).** Gene prioritization was informed by the FDA-recognized content in OncoKB [44].

Cancer	Gene list
Skin	CDK6;FANCE;JAK3;NRG1;RUNX1;ACVR1;CDK8;FANCF;JUN;NSD1;RYBP;ACVR1B;CDKN1A;FANCG;KDM5A;NTRK1;SDC4;AKT1; CDKN1B;FANCL;KDM5C;NTRK2;SDHA;AKT2;CDKN2A;FAT1;KDM6A;NTRK3;SDHAF2;AKT3;CDKN2B;FBXW7;KDR;NUTM1;SDHB; ALK;CDKN2C;FGF10;KEAP1;PAK7;SDHC;ALOX12B;CEBPA;FGF14;KEL;PALB2;SDHD;ANKRD11;CHD4;FGF19;KIF5B;PARK2;SETD2; APC;CHEK1;FGF23;KIT;PARP1;SF3B1;AR;CHEK2;FGF3;KLF4;PAX5;SGK1;ARAF;CIC;FGF4;KLHL6;PBRM1;SH2B3;ARFRP1;CREBBP; FGF6;KMT2A;PDCD1;SLC34A2;ARID1A;CRKL;FGFR1;KMT2B;PDCD1LG2;SMAD2;ARID1B;CRLF2;FGFR2;KMT2C;PDGFRA;SMAD3; ARID2;CSF1R;FGFR3;KMT2D;PDGFRB;SMAD4;ARID5B;CSF3R;FGFR4;KRAS;PDK1;SMARCA4;ASXL1;CTCF;FH;LATS1;PIK3C2B; SMARCB1;ASXL2;CTLA4;FLCN;LATS2;PIK3C2G;SMO;ATIC;CTNNA1;FLI1;LMO1;PIK3CA;SNCAIP;ATM;CTNNB1;FLT1;LRP1B; PIK3CB;SOCS1;ATR;CUL3;FLT3;LYN;PIK3CD;SOX17;ATRX;CUX1;FLT4;MAP2K1;PIK3R1;SOX2;AURKA;CXCR4;FOXA1;MAP2K2; PIK3R2;SOX9;AURKB;CYP2C8;FOXL2;MAP2K4;PIM1;SPEN;AXIN1;CYP2D6;FOXO1;MAP3K1;PMS1;SPOP;AXIN2;DAXX;FOXP1; MAP3K13;PMS2;SRC;AXL;DCUN1D1;FUBP1;MAPK1;POLD1;STAG2;B2M;DDR2;GABRA6;MAPK3;POLE;STAT3;BAP1;DHFR;GATA3; MAX;PPARG;STAT5B;BARD1;DICER1;GATA4;MCL1;PPM1D;STK11;BBC3;DNMT1;GATA6;MDM2;PPP2R1A;SUFU;BCL10;DNMT3A; GID4;MDM4;PPP2R2A;SYK;BCL2;DNMT3B;GLI1;MED12;PPP6C;TBX3;BCL2L1;DOT1L;GNA11;MEF2B;PRDM1;TCF3;BCL2L11;DPYD; GNA13;MEN1;PREX2;TCF7L2;BCL2L2;DYNC2H1;GNAQ;MET;PRKAR1A;TEK;BCL6;EED;GNAS;MGA;PRKCI;TERT;BCOR;EGFR; GPS2;MITF;PTCH1;TET1;BCORL1;EIF1AX;GRIN2A;MLH1;PTEN;TET2;BCR;EML4;GRM3;MPL;PTPN11;TGFBR2;BLM;EMSY;GSK3B; MRE11A;PTPRD;TMPRSS2;BMPR1A;EP300;GSTM1;MSH2;PTPRS;TNFAIP3;BRAF;EPCAM;GSTP1;MSH3;PTPRT;TNFRSF14;BRCA1; EPHA3;GSTT1;MSH6;QKI;TOP1;BRCA2;EPHA5;H3F3A;MST1;RAC1;TP53;BRD4;EPHA7;HDAC1;MST1R;RAD21;TP63;BRIP1;EPHB1; HGF;MTHFR;RAD50;TRAF7;BTG1;ERBB2;HIST1H3B;MTOR;RAD51;TSC1;BTK;ERBB3;HLA-A;MUTYH;RAD51B;TSC2;CALR;ERBB4; HLA-B;MYB;RAD51C;TSHR;CARD11;ERCC1;HNF1A;MYC;RAD51D;U2AF1;CASP8;ERCC2;HOXB13;MYCN;RAD52;UGT1A1;CBFB; ERCC3;HRAS;MYD88;RAD54L;VEGFA;CBL;ERCC4;HSD3B1;MYOD1;RAF1;VHL;CBR3;ERCC5;ID3;NBN;RANBP2;VTCN1;CCND1; ERG;IDH1;NCOR1;RARA;WHSC1;CCND2;ERRFI1;IDH2;NF1;RASA1;WHSC1L1;CCND3;ESR1;IGF1;NF2;RB1;WT1;CCNE1;ETV1; IGF1R;NFE2L2;RBM10;XPC;CD274;ETV4;IGF2;NFKBIA;RECQL4;XPO1;CD276;ETV5;IKBKE;NKX3-1;REL;XRCC1;CD74;ETV6;IKZF1; NOTCH;RET;XRCC2;CD79A;EWSR1;INHA;NOTCH1;RHOA;ZNF217;CD79B;EZH2;INPP4B;NOTCH2;RICTOR;ZNF703;CDA;EZR;IRF2; NOTCH3;RIT1;CDC73;FAM175A;IRF4;NOTCH4;RNF43;CDH1;FAM46C;IRS2;NPM1;ROS1;CDK12;FANCA;JAK1;NQO1;RPTOR;CDK4; FANCC;JAK2;NRAS;RSPO2;
Liver	CCND1;FGF19;MDM2;PALB2;TP53;ALK;CDA;FGF3;MET;PBRM1;TSC1;APC;CDKN2A;FGF4;MLH1;PIK3CA;TSC2;ARID1A;CDKN2B; FGFR1;MSH2;PMS2;UGT1A1;ARID2;CTNNB1;FGFR2;MSH6;POLD1;VEGFA;ATIC;CYP2C8;FGFR3;MTHFR;POLE;XPC;ATM;DHFR; GSTP1;MTOR;PTEN;XRCC1;ATR;DPYD;IDH1;MYC;RAD51D;BAP1;DYNC2H1;IDH2;NF1;RB1;BRAF;EGFR;JAK1;NRAS;RET;BRCA1; EPCAM;KEAP1;NTRK1;ROS1;BRCA2;ERBB2;KRAS;NTRK2;SMAD4;CBR3;ERCC1;MAP3K1;NTRK3;TERT;
Prostate	BRIP1;CYP2C8;HOXB13;NTRK1;RAD54L;ALK;CBR3;DHFR;HRAS;NTRK2;RAF1;APC;CCND1;DPYD;IDH1;NTRK3;RB1;AR;CDA; DYNC2H1;KRAS;PALB2;RET;ATIC;CDH1;EGFR;MAP3K1;PIK3CA;SPOP;ATM;CDK12;EPCAM;MED12;PMS2;TMPRSS2;ATR;CDK4; ERBB2;MLH1;POLE;TP53;ATRX;CDK6;ERCC1;MRE11A;PPP2R2A;UGT1A1;BAP1;CDKN2A;FAM175A;MSH2;PTEN;XPC;BARD1; CDKN2B;FANCA;MSH6;RAD50;XRCC1;BRAF;CHEK1;FANCL;MTHFR;RAD51B;BRCA1;CHEK2;FOXA1;MYC;RAD51C;BRCA2; CTNNB1;GSTP1;NBN;RAD51D;
Bladder	CDA;ERBB2;MAP3K1;NTRK2;RET;APC;CDK12;ERCC1;MDM2;NTRK3;ROS1;ARID1A;CDK4;FANCA;MET;PALB2;SDHA;ARID2;CDK6; FANCL;MLH1;PBRM1;SETD2;ATIC;CDKN2A;FBXW7;MRE11A;PIK3CA;SMARCA4;ATM;CDKN2B;FGF19;MSH2;PMS2;SOX2;ATR; CHEK1;FGF3;MSH6;POLE;STK11;ATRX;CHEK2;FGF4;MTHFR;PPP2R2A;TERT;BAP1;CREBBP;FGFR1;MTOR;PTEN;TP53;BARD1; CTNNB1;FGFR2;MUTYH;RAD50;TSC1;BRAF;CYP2C8;FGFR3;MYC;RAD51B;TSC2;BRCA1;DHFR;GATA3;NBN;RAD51C;UGT1A1; BRCA2;DPYD;GSTP1;NF1;RAD51D;XPC;BRIP1;DYNC2H1;HRAS;NOTCH2;RAD54L;XRCC1;CBR3;EGFR;KRAS;NRAS;RAF1;CCND1; EPCAM;MAP2K1;NTRK1;RB1;

Supplementary Table 3: **ST prediction performance of different foundation models on HEST-1k dataset (Part 1/2).** The 95% CI is included in parentheses. Best model is **bolded**, second-best is underlined.

Model	Organ	PCC	Num	Test patient num
UNI 2	breast	0.160 (0.088–0.231)	125	36
Virchow2	breast	0.135 (0.065–0.205)	125	36
H-optimus-1	breast	0.145 (0.097–0.194)	125	36
Prov-Gigapath	breast	0.121 (0.056–0.187)	125	36
Phikon-v2	breast	0.136 (0.056–0.217)	125	36
Shazam	breast	**0.176 (0.076–0.276)**	125	36
UNI 2	prostate	0.138 (-0.082–0.359)	60	2
Virchow2	prostate	0.138 (-0.177–0.453)	60	2
Prov-Gigapath	prostate	0.133 (-0.059–0.325)	60	2
Phikon-v2	prostate	0.141 (-0.281–0.562)	60	2
H-optimus-1	prostate	0.156 (-0.223–0.535)	60	2
Shazam	prostate	**0.185 (-0.181–0.550)**	60	2
UNI 2	skin	0.103 (0.048–0.158)	27	7
Virchow2	skin	0.108 (0.048–0.167)	27	7
Prov-Gigapath	skin	0.099 (0.035–0.162)	27	7
Phikon-v2	skin	0.100 (0.046–0.154)	27	7
H-optimus-1	skin	0.108 (0.053–0.163)	27	7
Shazam	skin	**0.153 (0.072–0.234)**	27	7
UNI 2	bowel	0.193 (0.126–0.259)	64	6
Virchow2	bowel	0.186 (0.137–0.236)	64	6
H-optimus-1	bowel	0.165 (0.106–0.225)	64	6
Prov-Gigapath	bowel	0.132 (0.008–0.256)	64	6
Phikon-v2	bowel	0.156 (0.075–0.237)	64	6
Shazam	bowel	**0.242 (0.139–0.345)**	64	6

Supplementary Table 4: **ST prediction performance of different foundation models on HEST-1k dataset (Part 2/2).** The 95% CI is included in parentheses. Best model is **bolded**, second-best is underlined.

Model	Organ	PCC	Num	Test patient num
UNI 2	liver	0.050 (0.048–0.051)	13	2
Virchow2	liver	0.070 (-0.121–0.262)	13	2
Prov-Gigapath	liver	0.060 (-0.159–0.278)	13	2
Phikon-v2	liver	0.059 (0.021–0.097)	13	2
H-optimus-1	liver	0.057 (-0.037–0.152)	13	2
Shazam	liver	**0.145 (-0.015–0.304)**	13	2
UNI 2	kidney	0.210 (0.161–0.259)	25	24
Virchow2	kidney	0.201 (0.156–0.246)	25	24
H-optimus-1	kidney	0.201 (0.167–0.235)	25	24
Prov-Gigapath	kidney	0.207 (0.184–0.230)	25	24
Phikon-v2	kidney	0.217 (0.190–0.244)	25	24
Shazam	kidney	**0.236 (0.187–0.285)**	25	24
UNI 2	brain	0.086 (0.060–0.112)	24	11
Virchow2	brain	0.083 (0.067–0.100)	24	11
Prov-Gigapath	brain	0.086 (0.068–0.104)	24	11
Phikon-v2	brain	0.092 (0.074–0.110)	24	11
H-optimus-1	brain	0.093 (0.067–0.119)	24	11
Shazam	brain	**0.127 (0.097–0.158)**	24	11
UNI 2	bladder	0.049 (-0.003–0.101)	6	5
Virchow2	bladder	0.061 (0.006–0.115)	6	5
H-optimus-1	bladder	0.033 (-0.026–0.092)	6	5
Prov-Gigapath	bladder	0.050 (0.005–0.095)	6	5
Phikon-v2	bladder	0.047 (-0.017–0.112)	6	5
Shazam	bladder	**0.107 (0.020–0.194)**	6	5

Supplementary Table 6: **Performance of Survival Analysis on TCGA-COADREAD, TCGA-GBMLGG, and TCGA-KIRC Datasets.** The 95% CI is included in parentheses. The best and second-best performed models are **bolded** and underlined.

Model	TCGA-COADREAD	TCGA-GBMLGG	TCGA-KIRC
UNI 2	0.596 (0.540–0.652)	0.756 (0.715–0.797)	0.687 (0.605–0.769)
Virchow2	0.625 (0.558–0.691)	0.756 (0.708–0.805)	0.694 (0.638–0.750)
H-optimus-1	0.646 (0.559–0.732)	0.767 (0.729–0.804)	0.710 (0.640–0.780)
Prov-Gigapath	0.595 (0.542–0.648)	0.766 (0.724–0.809)	0.689 (0.631–0.747)
Phikon-v2	0.574 (0.516–0.633)	0.748 (0.715–0.782)	0.662 (0.584–0.741)
Shazam	**0.710 (0.655–0.765)**	**0.768 (0.726–0.811)**	**0.742 (0.683–0.802)**

Supplementary Table 7: **Performance of Survival Analysis on TCGA-LUAD, TCGA-LUSC, TCGA-SKCM, and TCGA-STAD Datasets.** The 95% CI is included in parentheses. The best and second-best performed models are **bolded** and underlined.

Model	TCGA-LUAD	TCGA-LUSC	TCGA-SKCM	TCGA-STAD
UNI 2	0.596 (0.493–0.700)	0.521 (0.433–0.609)	0.622 (0.589–0.655)	0.584 (0.512–0.657)
Virchow2	0.612 (0.520–0.705)	0.538 (0.456–0.619)	0.629 (0.590–0.668)	0.605 (0.564–0.647)
H-optimus-1	0.588 (0.476–0.699)	0.545 (0.459–0.630)	0.637 (0.613–0.662)	0.601 (0.554–0.648)
Prov-Gigapath	0.603 (0.474–0.731)	0.525 (0.450–0.599)	0.628 (0.604–0.653)	0.563 (0.517–0.608)
Phikon-v2	0.571 (0.507–0.635)	0.498 (0.416–0.579)	0.581 (0.537–0.625)	0.591 (0.515–0.668)
Shazam	**0.628 (0.577–0.679)**	**0.583 (0.525–0.641)**	**0.657 (0.613–0.701)**	**0.626 (0.531–0.721)**

Supplementary Table 8: **Tile-level classification performance of different foundation models on HunCRC dataset.** Non-parametric bootstrapping with 1,000 replicates is employed for statistical analysis. The 95% CI is included in parentheses. Best performing model is **bolded** and second-best is underlined.

Model	Balanced ACC	Weighted F1	Top1-ACC
UNI 2	0.520 (0.492–0.546)	0.439 (0.421–0.458)	0.812 (0.807–0.817)
Virchow2	**0.567 (0.545–0.588)**	0.443 (0.427–0.460)	0.797 (0.792–0.802)
H-optimus-1	0.458 (0.435–0.483)	0.414 (0.397–0.430)	0.829 (0.824–0.833)
Prov-Gigapath	0.472 (0.446–0.499)	0.428 (0.408–0.448)	**0.833 (0.828–0.838)**
Phikon-v2	0.472 (0.446–0.498)	0.409 (0.393–0.426)	0.814 (0.809–0.819)
Shazam	0.549 (0.523–0.575)	**0.459 (0.442–0.475)**	0.829 (0.824–0.834)

Supplementary Table 9: **Tile-level classification performance of different foundation models on UniToPatho dataset.** Non-parametric bootstrapping with 1,000 replicates is employed for statistical analysis. The 95% CI is included in parentheses. Best performing model is **bolded** and second-best is underlined.

Model	Balanced ACC	Weighted F1	Top1-ACC
UNI 2	0.560 (0.539–0.583)	0.536 (0.515–0.558)	0.560 (0.543–0.579)
Virchow2	**0.606 (0.584–0.628)**	**0.586 (0.565–0.608)**	0.610 (0.590–0.630)
H-optimus-1	0.465 (0.444–0.486)	0.423 (0.401–0.443)	0.505 (0.486–0.525)
Prov-Gigapath	0.498 (0.479–0.518)	0.508 (0.488–0.529)	0.595 (0.576–0.616)
Phikon-v2	0.540 (0.519–0.562)	0.497 (0.474–0.518)	0.553 (0.534–0.572)
Shazam	0.532 (0.511–0.553)	0.548 (0.525–0.570)	**0.622 (0.604–0.643)**

Supplementary Table 10: **Tile-level classification performance of different foundation models on PCAM dataset.** Non-parametric bootstrapping with 1,000 replicates is employed for statistical analysis. The 95% CI is included in parentheses. Best performing model is **bolded** and second-best is underlined.

Model	Balanced ACC	Weighted F1	Top1-ACC
UNI 2	0.932 (0.929–0.934)	0.932 (0.929–0.934)	0.932 (0.929–0.935)
Virchow2	0.929 (0.927–0.932)	0.929 (0.926–0.932)	0.929 (0.927–0.932)
H-optimus-1	0.935 (0.932–0.938)	0.935 (0.932–0.937)	0.935 (0.932–0.938)
Prov-Gigapath	0.930 (0.927–0.932)	0.929 (0.927–0.932)	0.930 (0.927–0.932)
Phikon-v2	0.886 (0.882–0.889)	0.885 (0.881–0.888)	0.886 (0.882–0.889)
Shazam	**0.944 (0.942–0.947)**	**0.944 (0.941–0.947)**	**0.944 (0.941–0.947)**

Supplementary Table 11: **Tile-level classification performance of different foundation models on ESCA dataset.** Non-parametric bootstrapping with 1,000 replicates is employed for statistical analysis. The 95% CI is included in parentheses. Best performing model is **bolded** and second-best is underlined.

Model	Balanced ACC	Weighted F1	Top1-ACC
UNI 2	0.808 (0.803–0.813)	0.832 (0.827–0.836)	0.945 (0.944–0.946)
Virchow2	0.848 (0.843–0.853)	0.867 (0.863–0.871)	**0.956 (0.955–0.957)**
H-optimus-1	0.812 (0.807–0.817)	0.837 (0.833–0.842)	0.951 (0.950–0.952)
Prov-Gigapath	0.782 (0.778–0.787)	0.787 (0.782–0.791)	0.921 (0.919–0.922)
Phikon-v2	0.762 (0.757–0.767)	0.736 (0.730–0.740)	0.879 (0.878–0.881)
Shazam	**0.868 (0.863–0.873)**	**0.881 (0.878–0.885)**	0.953 (0.952–0.954)

Supplementary Table 12: **Tile-level classification performance of different foundation models on PanCancer-TIL dataset.** Non-parametric bootstrapping with 1,000 replicates is employed for statistical analysis. The 95% CI is included in parentheses. Best performing model is **bolded** and second-best is underlined.

Model	Balanced ACC	Weighted F1	Top1-ACC
UNI 2	**0.918 (0.915–0.921)**	0.908 (0.905–0.911)	0.943 (0.941–0.945)
Virchow2	0.902 (0.898–0.906)	0.910 (0.907–0.913)	**0.946 (0.945–0.948)**
H-optimus-1	0.912 (0.908–0.915)	0.902 (0.899–0.905)	0.939 (0.937–0.941)
Prov-Gigapath	0.908 (0.904–0.911)	0.898 (0.895–0.901)	0.937 (0.935–0.938)
Phikon-v2	0.913 (0.910–0.916)	0.909 (0.906–0.912)	0.944 (0.942–0.946)
Shazam	0.916 (0.912–0.919)	**0.912 (0.909–0.915)**	0.946 (0.944–0.948)

Supplementary Table 13: **Tile-level classification performance of different foundation models on CRC-100k dataset.** Non-parametric bootstrapping with 1,000 replicates is employed for statistical analysis. The 95% CI is included in parentheses. Best performing model is **bolded** and second-best is underlined.

Model	Balanced ACC	Weighted F1	Top1-ACC
UNI 2	0.953 (0.947–0.958)	0.949 (0.943–0.954)	0.960 (0.956–0.964)
Virchow2	0.953 (0.947–0.958)	0.948 (0.941–0.954)	0.967 (0.962–0.971)
H-optimus-1	0.952 (0.946–0.957)	0.946 (0.939–0.953)	0.965 (0.961–0.970)
Prov-Gigapath	0.952 (0.946–0.957)	0.948 (0.943–0.954)	0.964 (0.959–0.969)
Phikon-v2	0.928 (0.922–0.935)	0.925 (0.918–0.932)	0.950 (0.945–0.954)
Shazam	**0.960 (0.955–0.966)**	**0.956 (0.949–0.962)**	**0.972 (0.968–0.977)**

Supplementary Table 14: **Tile-level classification performance of different foundation models on CCRCC dataset.** Non-parametric bootstrapping with 1,000 replicates is employed for statistical analysis. The 95% CI is included in parentheses. Best performing model is **bolded** and second-best is underlined.

Model	Balanced ACC	Weighted F1	Top1-ACC
UNI 2	0.964 (0.954–0.972)	0.960 (0.952–0.967)	0.974 (0.970–0.979)
Virchow2	0.964 (0.955–0.972)	0.961 (0.952–0.968)	0.972 (0.968–0.976)
H-optimus-1	0.960 (0.951–0.968)	0.957 (0.949–0.965)	0.972 (0.968–0.976)
Prov-Gigapath	0.958 (0.949–0.967)	0.959 (0.951–0.966)	0.971 (0.966–0.975)
Phikon-v2	**0.965 (0.957–0.973)**	0.959 (0.950–0.967)	0.973 (0.968–0.977)
Shazam	0.963 (0.953–0.971)	**0.975 (0.971–0.978)**	**0.975 (0.971–0.979)**

Supplementary Table 15: **Tile-level classification performance of different foundation models on BACH dataset.** Non-parametric bootstrapping with 1,000 replicates is employed for statistical analysis. The 95% CI is included in parentheses. Best performing model is **bolded** and second-best is underlined.

Model	Balanced ACC	Weighted F1	Top1-ACC
UNI 2	**0.988 (0.958–1.000)**	**0.987 (0.958–1.000)**	**0.988 (0.963–1.000)**
Virchow2	**0.988 (0.958–1.000)**	**0.987 (0.958–1.000)**	**0.988 (0.963–1.000)**
H-optimus-1	0.963 (0.916–1.000)	0.962 (0.913–1.000)	0.963 (0.912–1.000)
Prov-Gigapath	0.912 (0.844–0.969)	0.915 (0.848–0.972)	0.912 (0.850–0.975)
Phikon-v2	0.825 (0.738–0.901)	0.829 (0.741–0.905)	0.825 (0.738–0.900)
Shazam	**0.988 (0.958–1.000)**	**0.987 (0.958–1.000)**	**0.988 (0.963–1.000)**

Supplementary Table 16: **Tile-level classification performance of different foundation models on CRC-MSI dataset.** Non-parametric bootstrapping with 1,000 replicates is employed for statistical analysis. The 95% CI is included in parentheses. Best performing model is **bolded** and second-best is underlined.

Model	Balanced ACC	Weighted F1	Top1-ACC
UNI 2	0.736 (0.729–0.743)	0.713 (0.707–0.719)	0.829 (0.826–0.834)
Virchow2	0.733 (0.726–0.740)	0.722 (0.715–0.728)	0.842 (0.838–0.846)
H-optimus-1	0.721 (0.714–0.728)	0.700 (0.694–0.706)	0.822 (0.818–0.826)
Prov-Gigapath	0.730 (0.723–0.737)	0.693 (0.687–0.699)	0.806 (0.802–0.810)
Phikon-v2	0.633 (0.626–0.640)	0.633 (0.626–0.639)	0.801 (0.796–0.805)
Shazam	**0.738 (0.731–0.745)**	**0.726 (0.720–0.732)**	**0.844 (0.840–0.848)**

Supplementary Table 17: **Tile-level classification performance of different foundation models on PanCancer-TCGA dataset.** Non-parametric bootstrapping with 1,000 replicates is employed for statistical analysis. The 95% CI is included in parentheses. Best performing model is **bolded** and second-best is underlined.

Model	Balanced ACC	Weighted F1	Top1-ACC
UNI 2	0.919 (0.916–0.922)	0.925 (0.922–0.927)	0.938 (0.936–0.940)
Virchow2	0.914 (0.911–0.917)	0.920 (0.917–0.923)	0.934 (0.933–0.936)
H-optimus-1	0.927 (0.924–0.930)	**0.932 (0.930–0.935)**	**0.946 (0.945–0.948)**
Prov-Gigapath	0.909 (0.906–0.912)	0.913 (0.910–0.916)	0.930 (0.929–0.932)
Phikon-v2	0.909 (0.905–0.912)	0.916 (0.913–0.918)	0.929 (0.928–0.931)
Shazam	**0.928 (0.925–0.930)**	0.931 (0.929–0.934)	0.945 (0.944–0.947)

Supplementary Table 21: **Complementarity of multi-level features for WSI-level survival prediction on TCGA-COADREAD.** The 95% CI is included in parentheses. The best and second-best performed models are **bolded** and underlined.

Model	C-index
Partial feature combinations	
Low only	0.680 (0.629–0.731)
Mid only	0.689 (0.636–0.743)
High only	0.686 (0.638–0.735)
Low + High	0.702 (0.618–0.785)
Low + Mid	0.692 (0.638–0.747)
Mid + High	0.699 (0.602–0.796)
Full multi-level feature ensemble	
Shazam (Low + Mid + High)	**0.710 (0.655–0.765)**

Supplementary Table 22: **Ablation study evaluating the Mixture of Experts (MoE) module in Shazam on the CRC-100K dataset.** Best performance is shown in **bold**, and the second-best is underlined.

Model	Balanced ACC	Weighted F1	Top1-ACC
Shazam	**0.960 (0.955–0.966)**	**0.956 (0.949–0.962)**	**0.972 (0.968–0.977)**
MoE Ablation			
w/o MoE	0.955 (0.950–0.960)	0.951 (0.946–0.956)	0.970 (0.967–0.974)

Supplementary Table 23: **Ablation study evaluating the Mixture of Experts (MoE) module in Shazam on the CCRCC dataset**. Best performance is shown in **bold**, and the second-best is underlined.

Model	Balanced ACC	Weighted F1	Top1-ACC
Shazam	**0.963 (0.953–0.971)**	**0.975 (0.971–0.978)**	**0.975 (0.971–0.979)**
MoE Ablation			
w/o MoE	**0.963 (0.954–0.971)**	0.973 (0.969–0.977)	0.973 (0.969–0.977)